LEARNING TO TRANSFER WITH TRIPLY ADVERSARIAL NETS

By Gilles Louppe

New York University

Transfer learning, and its variants known as learning under co-variate shift or domain adaptation, arises whenever test instances are governed by a distribution that may be arbitrarily different from the distribution of the training instances. This problem has traditionally been solved either by reweighting training instances or by learning robust feature representations over domains. In this work, we propose a new paradigm which consists instead in learning how to transform training instances into test instances, possibly across input spaces of distinct dimensions, structures or supports. For this purpose, we extend the generative adversarial networks framework of Goodfellow et al. (2014) to a triply adversarial process: a transformer network T for generating test instances from training instances, a discriminative network D for estimating whether an instance comes either from the training or the test distributions, and a classifier network $C \circ T$ for classifying training instances in the projected space. This 3-player game results in a network T capable of transforming training instances into test instances, while preserving the separation between classes, so that C should be an accurate classifier on test data. Experimental results demonstrate the potential of this novative approach for transfer learning. More fundamentally, this paradigm also raises interesting theoretical issues, since transformations may not always be unique nor grounded, depending on the studied problem.

1. **Introduction.**

2. **Method.**

3. **Experiments.**

4. **Related work.**

5. **Conclusions.**

References.

Goodfellow, I., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., Courville, A., and Bengio, Y. (2014). Generative adversarial nets. In *Advances in Neural Information Processing Systems*, pages 2672–2680.